

# The PageRank Algorithm (Chapter 5)

---

Wojtek Kowalczyk

*plus slides*

<http://www.mmds.org/mmds/v2.1/ch05-linkanalysis2.pptx>

# Some history

---

- ❑ Internet in early 80': small, slow, local,
- ❑ Internet in early 90': growing, global, slow
  - newsgroups
  - mailing lists
  - bulletin boards
- ❑ “primitive” search engines:
  - Lycos
  - Excite
  - Netscape
  - Yahoo!
  - MSN
  - “hybrid”: combinations of existing engines
- ❑ **1997: google.stanford.edu; 15 September 1997: domain google.com registered**

# The key to success: the PageRank algorithm

---

- ❑ Invented by Sergey Brin and Larry Page in 1996 (Stanford University)
- ❑ Estimates “page importance” by analyzing the structure of links
- ❑ Some Rough Statistics (from August 29th, 1996)
  - Total indexable HTML urls: 75 Million
  - Total content downloaded: 207 gigabytes ...
  - BackRub is written in Java and Python and runs on several Sun Ultras and Intel Pentiums running Linux. The primary database is kept on a Sun Ultra II with 28GB of disk. Scott Hassan and Alan Steremberg have provided a great deal of very talented implementation help. **Sergey Brin** has also been very involved and deserves many thanks.  
-**Larry Page** [pagescs.stanford.edu](http://pagescs.stanford.edu)
- ❑ By the end of 1998, Google had an index of about 60 million pages
- ❑ [http://en.wikipedia.org/wiki/History\\_of\\_Google](http://en.wikipedia.org/wiki/History_of_Google)

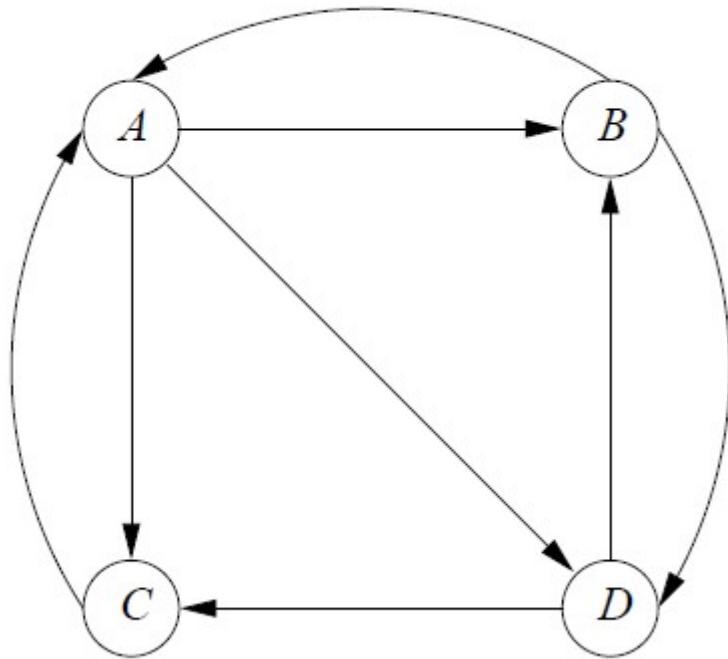
# The PageRank Algorithm (Ch. 5)

---

- ❑ How Google determines the importance of a page?
- ❑ A “random surfer” model:
  - visitors start their session at random pages
  - visitors walk along links at random
  - choices are made uniformly (each outgoing link has the same chance)
  - “page importance” = “frequency the surfer visits the page”
- ❑ It can be modeled by a Markov Process
  - transition matrix
  - iterative calculation of page probability distribution
  - solution = principal eigen vector of the transition matrix
- ❑ Nowadays a much more sophisticated model is used ...

# The transition matrix

---



$$M = \begin{matrix} & \begin{matrix} A & B & C & D \end{matrix} \\ \begin{matrix} A \\ B \\ C \\ D \end{matrix} & \begin{bmatrix} 0 & 1/2 & 1 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{bmatrix} \end{matrix}$$

Each column X lists probabilities of going to Y

$$v = [1/4, 1/4, 1/4, 1/4]'$$

What is  $Mv$ ? What is  $M(Mv)$ ?  $M(M(Mv))$ ?...

Entries in each column sum up to 1

## The transition matrix: iterating $Mv$

---

$$M = \begin{matrix} & \begin{matrix} A & B & C & D \end{matrix} \\ \begin{matrix} A \\ B \\ C \\ D \end{matrix} & \begin{bmatrix} 0 & 1/2 & 1 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{bmatrix} \end{matrix}$$


$$\begin{bmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{bmatrix} \begin{bmatrix} 9/24 \\ 5/24 \\ 5/24 \\ 5/24 \end{bmatrix} \begin{bmatrix} 15/48 \\ 11/48 \\ 11/48 \\ 11/48 \end{bmatrix} \begin{bmatrix} 11/32 \\ 7/32 \\ 7/32 \\ 7/32 \end{bmatrix} \cdots \begin{bmatrix} 3/9 \\ 2/9 \\ 2/9 \\ 2/9 \end{bmatrix}$$

# Convergence of the Markov process

---

- If:
  - the graph is strongly connected  
(i.e., there is a path between any two nodes)
  - there are no “dead ends” (nodes with no links going out)
- then
  - the sequence  $v, Mv, M(Mv), \dots$  converges to  $v'$  such that  $v' = Mv'$
  - $v'$  is the principal eigen vector of matrix  $M$   
(principal = biggest eigen value)
- In practice, 50-70 iterations are sufficient
- ... even for huge  $M$  (billions x billions , sparse) ...
- For small  $M$ ,  $v = Mv$  can be solved directly...

# Dead-ends

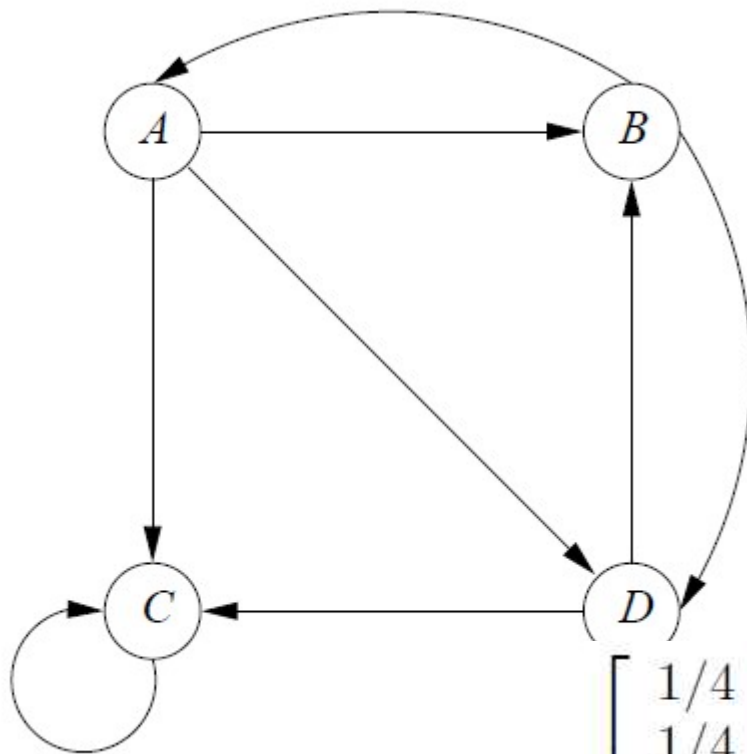


$$M = \begin{matrix} & \begin{matrix} A & B & C & D \end{matrix} \\ \begin{matrix} A \\ B \\ C \\ D \end{matrix} & \begin{bmatrix} 0 & 1/2 & 0 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{bmatrix} \end{matrix}$$

$$\begin{bmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{bmatrix} \begin{bmatrix} 3/24 \\ 5/24 \\ 5/24 \\ 5/24 \end{bmatrix} \begin{bmatrix} 5/48 \\ 7/48 \\ 7/48 \\ 7/48 \end{bmatrix} \begin{bmatrix} 21/288 \\ 31/288 \\ 31/288 \\ 31/288 \end{bmatrix} \cdots \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$



# Spider-traps



$$M = \begin{bmatrix} 0 & 1/2 & 0 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 1 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{bmatrix} \begin{matrix} A \\ B \\ C \\ D \end{matrix}$$

$$\begin{bmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{bmatrix} \begin{bmatrix} 3/24 \\ 5/24 \\ 11/24 \\ 5/24 \end{bmatrix} \begin{bmatrix} 5/48 \\ 7/48 \\ 29/48 \\ 7/48 \end{bmatrix} \begin{bmatrix} 21/288 \\ 31/288 \\ 205/288 \\ 31/288 \end{bmatrix} \dots \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$$

# Teleporting

---

□ At each step, decide:

- with probability  $\beta$  ( $0 < \beta \leq 1$ ) continue random walk
- with probability  $(1 - \beta)$  jump to any other node

□ The corresponding update rule:

$$\mathbf{v}' = \beta \mathbf{M} \mathbf{v} + (1 - \beta) \mathbf{e} / n \quad (\mathbf{e} \text{ is a vector of } n \text{ 1's})$$

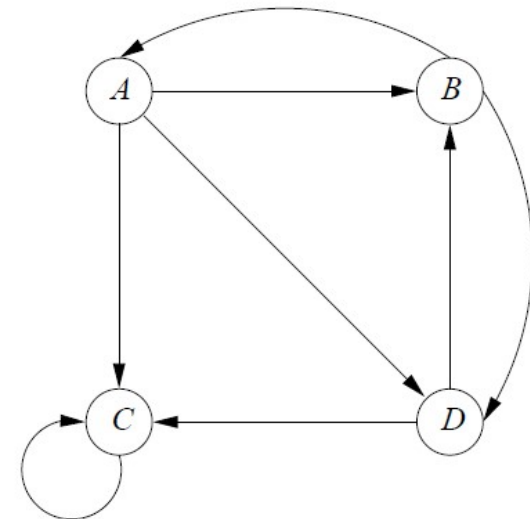
□  $\beta$  is usually 0.8 or 0.9

□ “adding extra links”: all problems solved (really?)

## Teleporting: example (slide 9 continued...)

$$\beta = 0.8 = 4/5$$

$$\mathbf{v}' = \begin{bmatrix} 0 & 2/5 & 0 & 0 \\ 4/15 & 0 & 0 & 2/5 \\ 4/15 & 0 & 4/5 & 2/5 \\ 4/15 & 2/5 & 0 & 0 \end{bmatrix} \mathbf{v} + \begin{bmatrix} 1/20 \\ 1/20 \\ 1/20 \\ 1/20 \end{bmatrix}$$



$$\begin{bmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{bmatrix} \begin{bmatrix} 9/60 \\ 13/60 \\ 25/60 \\ 13/60 \end{bmatrix} \begin{bmatrix} 41/300 \\ 53/300 \\ 153/300 \\ 53/300 \end{bmatrix} \begin{bmatrix} 543/4500 \\ 707/4500 \\ 2543/4500 \\ 707/4500 \end{bmatrix} \dots \begin{bmatrix} 15/148 \\ 19/148 \\ 95/148 \\ 19/148 \end{bmatrix}$$

## Teleporting: dead ends (slide 8 continued...)

---

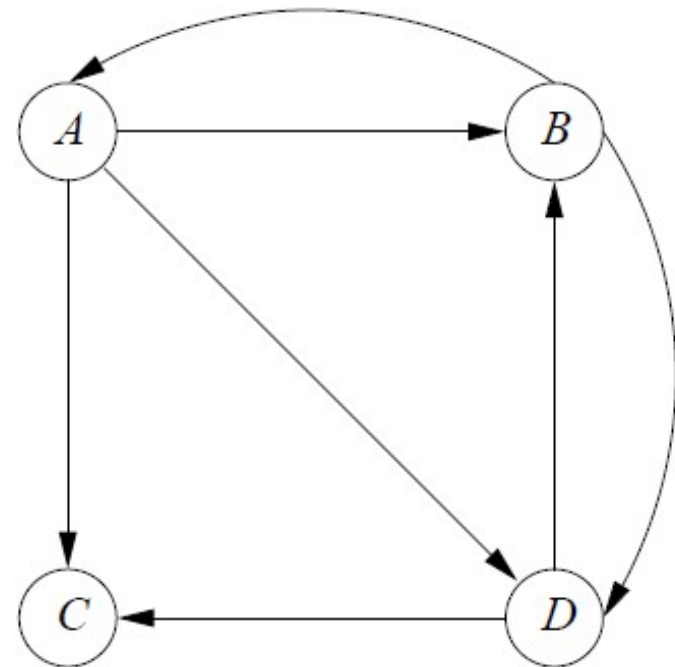
Beta=0.8

$$M = \begin{bmatrix} 0 & 1/2 & 0 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{bmatrix}$$
$$v = [1/4, 1/4, 1/4, 1/4]'$$

for i=1:20

$$v = \text{Beta} * M * v + (1 - \text{Beta}) * \text{ones}(\text{size}(v)) / \text{length}(v)$$

end



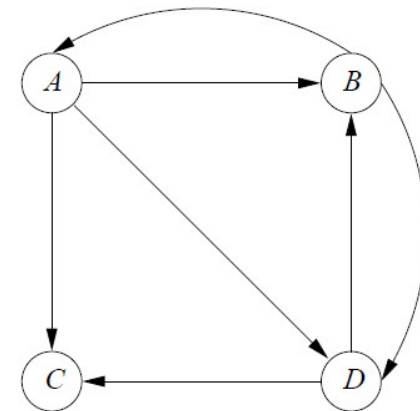
## Teleporting: dead ends (slide 8 continued...)

Convergence after 20 iterations:

Iter:	0	1	2	3	10	20
	0.2500	0.1500	0.1367	0.1207	0.1018	0.1014
	0.2500	0.2167	0.1767	0.1571	0.1290	0.1284
	0.2500	0.2167	0.1767	0.1571	0.1290	0.1284
	0.2500	0.2167	0.1767	0.1571	0.1290	0.1284

Does it make sense?

- **Why is  $\sum(v) < 1$ ?** Surfers “die” at dead ends? Not exactly...
- Is  $\sum(v)$  always  $> 0$  ? Yes: it is at least  $(1-\text{Beta})$
- What should be the case:
  - $p(A)=0.5*p(B)$ ?
  - $p(A)= 0.5*p(B) + (1-\text{Beta})/4$ ?
  - **$p(A)= \text{Beta}*0.5*p(B) + (1-\text{Beta})/4$ ? Check it!**



# Teleporting: dead-ends and spider traps

---

For graphs with no dead-ends “teleporting” works fine:  
the process converges to a vector  $v$ , such that  $\text{sum}(v)=1$   
and can be interpreted as “probabilities”.

For graphs with dead-ends, the “teleporting” still works,  
but the  $\text{sum}(v)$  may be smaller than 1 (no probabilistic interpretation).  
Still, useful in practice (used in the original “Google” paper).

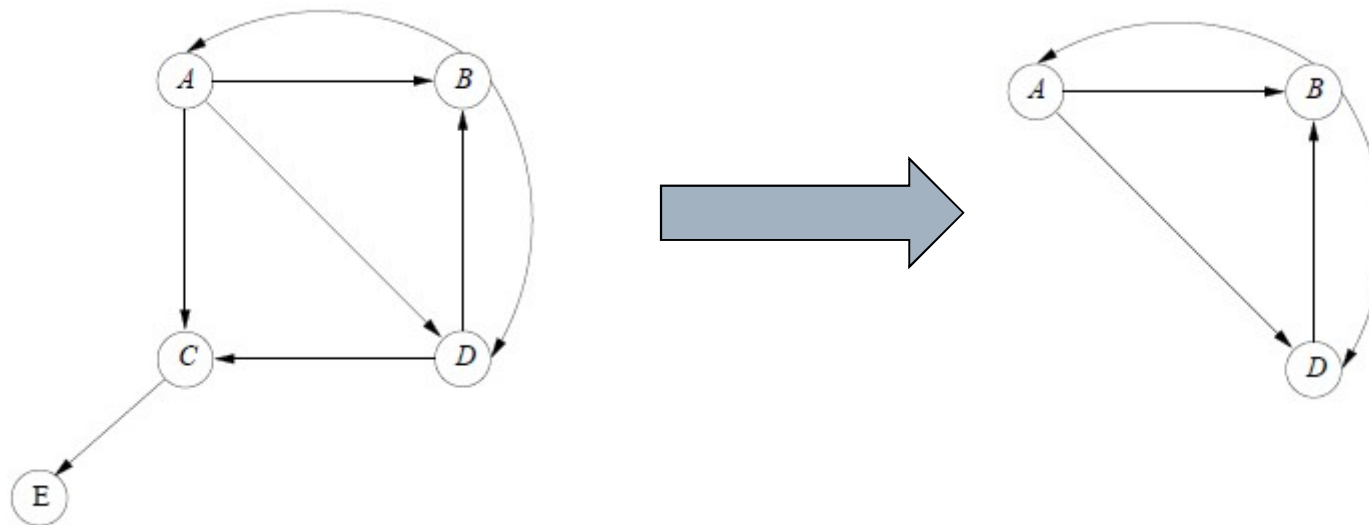
An alternative algorithm for handling dead-ends (read **Section 5.1.4!**)

- recursively remove dead ends and corresponding links,
- calculate of PageRank for the nodes of the remaining fragment of the graph,
- propagate computed values to the removed nodes.

Another alternative: when reaching a dead end jump the next node  
“uniformly at random with probability 1”

## Example from 5.1.4

---



After removal of nodes E and C, calculate PageRank of A, B, D and set:

$$\text{PageRank}(C) = 1/2 * \text{PageRank}(A) + 1/2 * \text{PageRank}(D)$$

$$\text{PageRank}(E) = \text{PageRank}(C)$$

Note that sum of PageRanks is now bigger than 1!

# Implementation of PageRank

---

## □ Key ideas:

- M is very sparse: say 10 links per page => 10 non-zeros in a column
- Use “inverted indexing” to represent M: a list of <node, outdegree, children of the node>
- keep on harddisk M and vector v\_old
- keep in RAM only v\_new
- update v\_new in a single scan of M and v\_old

## □ Some numbers:

- n=1.000.000.000 (1 billion nodes)
- RAM needed: 4GB (32bits per node)
- harddisk: about 40GB (10xRAM)
- a single scan: about 2-3 minutes
- 50 iterations => 2-3 hours



# Sparse Matrix Encoding

---

## □ Encode sparse matrix using only nonzero entries

- Space proportional roughly to number of links
- Say  $10N$ , or  $4 \times 10^1$  billion = 40GB
- **Still won't fit in memory, but will fit on disk**

source node	degree	destination nodes
0	3	1, 5, 7
1	5	17, 64, 113, 117, 245
2	2	13, 23

□ Assume enough RAM to fit  $\mathbf{v}^{new}$  into memory

■ Store  $\mathbf{v}^{old}$  and matrix  $\mathbf{M}$  on disk

□ 1 step of power-iteration is:

**Initialize** all entries of  $\mathbf{v}^{new} = (1-\beta) / N$

For each page  $i$  (of out-degree  $d_i$ ):

Read into memory:  $i, d_i, \text{dest}_1, \dots, \text{dest}_{d_i}, \mathbf{v}^{old}(i)$

For  $j = 1 \dots d_i$

$\mathbf{v}^{new}(\text{dest}_j) += \beta \mathbf{v}^{old}(i) / d_i$

0	
1	
2	
3	
4	
5	
6	

$\mathbf{v}^{new}$

source degree destination

0	3	1, 5, 6
1	4	17, 64, 113, 117
2	2	13, 23

$\mathbf{v}^{old}$	0
	1
	2
	3
	4
	5
	6

18

# In short: "column-wise" computations

---

When multiplying **M** by **v** we organize computations  
**"by columns"**: "a single, linear scan through the harddisk"

- data is already organized in this way,
- we access elements of "**old v**" one by one (a, b, c),
- outdegrees (one per column) are easy to find.

$$\begin{bmatrix} A & B & C \\ D & E & F \\ G & H & I \end{bmatrix} * \begin{bmatrix} a \\ b \\ c \end{bmatrix} = \begin{bmatrix} Aa + Bb + Cc \\ Da + Eb + Fc \\ Ga + Hb + Ic \end{bmatrix}$$

# A3: Structure of wikipedia links

---

Go to <https://zenodo.org/record/2539424> and fetch the file:

[https://zenodo.org/record/2539424/files/enwiki.wikilink\\_graph.2004-03-01.csv.gz?download=1](https://zenodo.org/record/2539424/files/enwiki.wikilink_graph.2004-03-01.csv.gz?download=1)

Investigate the graph:

- Dead ends
- Distribution of in-degrees
- Distribution of out-degrees of nodes
- Implement the page rank algorithm from slide 18
- Implement direct (sparse) matrix multiplication
- Compare results
- *Is this graph strongly connected?*

# Data preprocessing (prep)

---

The data has the following layout:

■	page_id_from	page_title_from	page_id_to	page_title_to
■	12	Anarchism	34568	16th century
■	12	Anarchism	35416	1793
■	...	...	...	...

- ☐ Extract only the 1<sup>st</sup> and the 3<sup>rd</sup> column
- ☐ Convert page\_id's into consecutive integers, in such a way that you can return back to the original numbering
- ☐ Both columns should use the same coding!
- ☐ Save the prepared data on HD

# Exploratory data analysis (eda)

---

- **Dead ends**: find nodes with no outgoing edges. How many have you found?
- **Distribution of in-degrees**: for every node compute the number of incoming edges
- **Distribution of out-degrees**: for every node compute the number of outgoing edges
- Make **nice & informative plots** of both distributions
- What is the **average out-degree** and the **average in-degree** of the graph?

## Estimate RAM requirements: (eda)

---

1. How much RAM would you need to store the transition matrix  $M$  and the initial vector  $v$  in RAM? Assume double precision (64 bits per number).
2. The same question assuming that you store  $M$  in a sparse matrix (in RAM)?
3. The same question, assuming that you use data structures as described on slide 17.

*embed your answers in the notebook [eda.ipynb](#)*

# Implement PageRank algorithm (*sparse*)

---

1. Store both  $M$  (as a sparse matrix) and  $v$  (in RAM).
2. Run 25 iterations of the “classical” update rule from slide 10, with  $\text{Beta}=0.8$ .
3. Plot the MSE of the differences (25 numbers):  $v - Mv$
4. Assume that your computer has 1GB RAM and the average out-degree of a graph  $G$  is 15.

*What is the maximal number of nodes of  $G$  such that your algorithm could be executed on your computer?*



## Implement PageRank algorithm from slide 18

---

1. Store both  $M$  and  $v^{old}$  and  $v^{new}$  in RAM.  
(Normally  $M$  and  $v^{old}$  stored on the hard disk)
2. Implement the algorithm from slide 18. Run 25 iterations of this algorithm with  $Beta=0.8$ .
3. Plot the MSE of the differences (25 numbers):  $v - Mv$
4. How much time is needed to run a single iteration?
5. Have you obtained similar/identical results as in the previous task? What might be the source of eventual differences?

# What to deliver?

---

Four Jupyter notebooks:

- `prep.ipynb`
- `eda.ipynb`
- `sparse.ipynb`
- `PageRank.ipynb`

that correspond to all the subtasks.

*Your notebooks should read/write files from/to the same directory as your notebooks. We will test your programs in a directory which contains `wikilink_graph.2004-03-01.csv`*